

AS6PR1

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
```

```
data = pd.read_csv("Breast_Cancer_Dataset.csv")
data.head()
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean
area_mean \					
0	842302	M	17.99	10.38	122.80
1001.0					
1	842517	M	20.57	17.77	132.90
1326.0					
2	84300903	M	19.69	21.25	130.00
1203.0					
3	84348301	M	11.42	20.38	77.58
386.1					
4	84358402	M	20.29	14.34	135.10
1297.0					

	smoothness_mean	compactness_mean	concavity_mean	concave
points_mean \				
0	0.11840	0.27760	0.3001	
0.14710				
1	0.08474	0.07864	0.0869	
0.07017				
2	0.10960	0.15990	0.1974	
0.12790				
3	0.14250	0.28390	0.2414	
0.10520				
4	0.10030	0.13280	0.1980	
0.10430				

	... texture_worst	perimeter_worst	area_worst	
smoothness_worst \				
0 ...	17.33	184.60	2019.0	0.1622

1	...	23.41	158.80	1956.0	0.1238
2	...	25.53	152.50	1709.0	0.1444
3	...	26.50	98.87	567.7	0.2098
4	...	16.67	152.20	1575.0	0.1374

	compactness_worst	concavity_worst	concave points_worst
0	0.6656	0.7119	0.2654
1	0.1866	0.2416	0.1860
2	0.4245	0.4504	0.2430
3	0.8663	0.6869	0.2575
4	0.2050	0.4000	0.1625

	fractal_dimension_worst	Unnamed: 32
0	0.11890	NaN
1	0.08902	NaN
2	0.08758	NaN
3	0.17300	NaN
4	0.07678	NaN

[5 rows x 33 columns]

data.isna().sum()

id	0
diagnosis	0
radius_mean	0
texture_mean	0
perimeter_mean	0
area_mean	0
smoothness_mean	0
compactness_mean	0
concavity_mean	0
concave points_mean	0
symmetry_mean	0
fractal_dimension_mean	0
radius_se	0
texture_se	0
perimeter_se	0
area_se	0
smoothness_se	0

```
compactness_se      0
concavity_se        0
concave points_se   0
symmetry_se         0
fractal_dimension_se 0
radius_worst        0
texture_worst       0
perimeter_worst     0
area_worst          0
smoothness_worst    0
compactness_worst   0
concavity_worst     0
concave points_worst 0
symmetry_worst      0
fractal_dimension_worst 0
Unnamed: 32         569
dtype: int64
```

```
data = data.drop(['Unnamed: 32'], axis=1)
```

```
le=LabelEncoder()
```

```
data["diagnosis"]= le.fit_transform(data["diagnosis"])
```

```
le.classes_
```

```
array([0, 1])
```

```
data.head()
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean
area_mean \					
0	842302	1	17.99	10.38	122.80
1001.0					
1	842517	1	20.57	17.77	132.90
1326.0					
2	84300903	1	19.69	21.25	130.00
1203.0					
3	84348301	1	11.42	20.38	77.58
386.1					
4	84358402	1	20.29	14.34	135.10
1297.0					

	smoothness_mean	compactness_mean	concavity_mean	concave
points_mean \				
0	0.11840	0.27760	0.3001	
0.14710				
1	0.08474	0.07864	0.0869	
0.07017				
2	0.10960	0.15990	0.1974	
0.12790				
3	0.14250	0.28390	0.2414	

```

0.10520
4          0.10030          0.13280          0.1980
0.10430

```

```

... radius_worst texture_worst perimeter_worst area_worst \
0 ...      25.38      17.33      184.60      2019.0
1 ...      24.99      23.41      158.80      1956.0
2 ...      23.57      25.53      152.50      1709.0
3 ...      14.91      26.50       98.87       567.7
4 ...      22.54      16.67      152.20      1575.0

```

```

smoothness_worst compactness_worst concavity_worst concave
points_worst \
0          0.1622          0.6656          0.7119
0.2654
1          0.1238          0.1866          0.2416
0.1860
2          0.1444          0.4245          0.4504
0.2430
3          0.2098          0.8663          0.6869
0.2575
4          0.1374          0.2050          0.4000
0.1625

```

```

symmetry_worst fractal_dimension_worst
0          0.4601          0.11890
1          0.2750          0.08902
2          0.3613          0.08758
3          0.6638          0.17300
4          0.2364          0.07678

```

```
[5 rows x 32 columns]
```

```

X = data.drop(columns = "diagnosis")
Y = data["diagnosis"]

```

```
X.head()
```

```

id radius_mean texture_mean perimeter_mean area_mean \
0 842302      17.99      10.38      122.80      1001.0
1 842517      20.57      17.77      132.90      1326.0
2 84300903     19.69      21.25      130.00      1203.0
3 84348301     11.42      20.38       77.58       386.1
4 84358402     20.29      14.34      135.10      1297.0

```

```

smoothness_mean compactness_mean concavity_mean concave
points_mean \
0          0.11840          0.27760          0.3001
0.14710
1          0.08474          0.07864          0.0869

```

```

0.07017
2      0.10960      0.15990      0.1974
0.12790
3      0.14250      0.28390      0.2414
0.10520
4      0.10030      0.13280      0.1980
0.10430

```

```

      symmetry_mean  ... radius_worst texture_worst perimeter_worst \
0      0.2419  ...      25.38      17.33      184.60
1      0.1812  ...      24.99      23.41      158.80
2      0.2069  ...      23.57      25.53      152.50
3      0.2597  ...      14.91      26.50      98.87
4      0.1809  ...      22.54      16.67      152.20

```

```

      area_worst smoothness_worst compactness_worst concavity_worst \
0      2019.0      0.1622      0.6656      0.7119
1      1956.0      0.1238      0.1866      0.2416
2      1709.0      0.1444      0.4245      0.4504
3      567.7      0.2098      0.8663      0.6869
4      1575.0      0.1374      0.2050      0.4000

```

```

      concave points_worst symmetry_worst fractal_dimension_worst
0      0.2654      0.4601      0.11890
1      0.1860      0.2750      0.08902
2      0.2430      0.3613      0.08758
3      0.2575      0.6638      0.17300
4      0.1625      0.2364      0.07678

```

```
[5 rows x 31 columns]
```

```
Y.head()
```

```

0      1
1      1
2      1
3      1
4      1

```

```
Name: diagnosis, dtype: int64
```

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size =
0.3,random_state = 1)
```

```
reg = LogisticRegression()
reg.fit(X_train,Y_train)
```

```
LogisticRegression()
```

```
Y_pred = reg.predict(X_test)
```

```

reg_acc = accuracy_score(Y_test,Y_pred)
print("Logisticregressionaccuracy=",reg_acc*100)

Logisticregressionaccuracy= 94.15204678362574

gaussian_nb=GaussianNB()
gaussian_nb.fit(X_train,Y_train)

GaussianNB()

Y_pred = gaussian_nb.predict(X_test)

bayes_acc = accuracy_score(Y_test,Y_pred)
print("Naive bayes accuracy=",bayes_acc*100)

Naive bayes accuracy= 94.15204678362574

knn5=KNeighborsClassifier(n_neighbors=5)
knn5.fit(X_train,Y_train)

KNeighborsClassifier()

Y_pred=knn5.predict(X_test)

KNN_acc = accuracy_score(Y_test,Y_pred)
KNN_acc*100

94.15204678362574

svc = SVC(C=3)
svc.fit(X_train,Y_train)

SVC(C=3)

Y_pred=svc.predict(X_test)
svc_acc=accuracy_score(Y_test,Y_pred)
svc_acc*100

63.1578947368421

dt=DecisionTreeClassifier(criterion='entropy',splitter='best',max_features=3,max_depth=3)
dt.fit(X_train,Y_train)

DecisionTreeClassifier(criterion='entropy', max_depth=3,
max_features=3)

Y_pred = dt.predict(X_test)
DT_acc = accuracy_score(Y_test,Y_pred)
print("Decision tree accuracy=",DT_acc*100)

Decision tree accuracy= 87.13450292397661

```

```

RF=RandomForestClassifier(n_estimators=100,
criterion='entropy',max_depth=3)
RF.fit(X_train,Y_train)

RandomForestClassifier(criterion='entropy', max_depth=3)

Y_pred = RF.predict(X_test)
RF_acc = accuracy_score(Y_test,Y_pred)
print("Random forest accuracy=",RF_acc*100)


Random forest accuracy= 94.15204678362574

con_mat = metrics.confusion_matrix(Y_test,Y_pred)

cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix =
con_mat, display_labels = [False, True])

cm_display.plot()
plt.show()

```



AS6PR2

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
```

```
data = pd.read_csv("Breast_Cancer_Dataset.csv")
```

```
data.isna().sum()
```

id	0
diagnosis	0
radius_mean	0
texture_mean	0
perimeter_mean	0
area_mean	0
smoothness_mean	0
compactness_mean	0
concavity_mean	0
concave points_mean	0
symmetry_mean	0
fractal_dimension_mean	0
radius_se	0
texture_se	0
perimeter_se	0
area_se	0
smoothness_se	0
compactness_se	0
concavity_se	0
concave points_se	0
symmetry_se	0
fractal_dimension_se	0


```

radius_worst      0
texture_worst     0
perimeter_worst   0
area_worst        0
smoothness_worst  0
compactness_worst 0
concavity_worst   0
concave points_worst 0
symmetry_worst    0
fractal_dimension_worst 0
Unnamed: 32       569
dtype: int64

```

```

data= data.drop(['Unnamed: 32'], axis=1)
le = LabelEncoder()
data["diagnosis"]= le.fit_transform(data["diagnosis"])
data

```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean \
0	842302	1	17.99	10.38	122.80
1	842517	1	20.57	17.77	132.90
2	84300903	1	19.69	21.25	130.00
3	84348301	1	11.42	20.38	77.58
4	84358402	1	20.29	14.34	135.10
...
564	926424	1	21.56	22.39	142.00
565	926682	1	20.13	28.25	131.20
566	926954	1	16.60	28.08	108.30
567	927241	1	20.60	29.33	140.10
568	92751	0	7.76	24.54	47.92

	area_mean	smoothness_mean	compactness_mean	concavity_mean \
0	1001.0	0.11840	0.27760	0.30010
1	1326.0	0.08474	0.07864	0.08690
2	1203.0	0.10960	0.15990	0.19740
3	386.1	0.14250	0.28390	0.24140
4	1297.0	0.10030	0.13280	0.19800
...
564	1479.0	0.11100	0.11590	0.24390
565	1261.0	0.09780	0.10340	0.14400
566	858.1	0.08455	0.10230	0.09251
567	1265.0	0.11780	0.27700	0.35140
568	181.0	0.05263	0.04362	0.00000

	concave points_mean	...	radius_worst	texture_worst
perimeter_worst \				
0	0.14710	...	25.380	17.33
184.60				

1	0.07017	...	24.990	23.41
158.80				
2	0.12790	...	23.570	25.53
152.50				
3	0.10520	...	14.910	26.50
98.87				
4	0.10430	...	22.540	16.67
152.20				
..
...				
564	0.13890	...	25.450	26.40
166.10				
565	0.09791	...	23.690	38.25
155.00				
566	0.05302	...	18.980	34.12
126.70				
567	0.15200	...	25.740	39.42
184.60				
568	0.00000	...	9.456	30.37
59.16				

	area_worst	smoothness_worst	compactness_worst	concavity_worst
\				
0	2019.0	0.16220	0.66560	0.7119
1	1956.0	0.12380	0.18660	0.2416
2	1709.0	0.14440	0.42450	0.4504
3	567.7	0.20980	0.86630	0.6869
4	1575.0	0.13740	0.20500	0.4000
..
564	2027.0	0.14100	0.21130	0.4107
565	1731.0	0.11660	0.19220	0.3215
566	1124.0	0.11390	0.30940	0.3403
567	1821.0	0.16500	0.86810	0.9387
568	268.6	0.08996	0.06444	0.0000

	concave	points_worst	symmetry_worst	fractal_dimension_worst
0		0.2654	0.4601	0.11890
1		0.1860	0.2750	0.08902

2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678
...
564	0.2216	0.2060	0.07115
565	0.1628	0.2572	0.06637
566	0.1418	0.2218	0.07820
567	0.2650	0.4087	0.12400
568	0.0000	0.2871	0.07039

[569 rows x 32 columns]

```
X = data.drop(columns = "diagnosis")
Y = data["diagnosis"]
```

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size =
0.3,random_state = 1)
```

```
reg = LogisticRegression()
reg.fit(X_train,Y_train)
Y_pred = reg.predict(X_test)
reg_acc = accuracy_score(Y_test,Y_pred)
```

```
gaussian_nb=GaussianNB()
gaussian_nb.fit(X_train,Y_train)
Y_pred = gaussian_nb.predict(X_test)
bayes_acc = accuracy_score(Y_test,Y_pred)
```

```
knn5=KNeighborsClassifier(n_neighbors=5)
knn5.fit(X_train,Y_train)
Y_pred=knn5.predict(X_test)
KNN_acc = accuracy_score(Y_test,Y_pred)
```

```
svc = SVC(C=10)
svc.fit(X_train,Y_train)
Y_pred=svc.predict(X_test)
svc_acc=accuracy_score(Y_test,Y_pred)
```

```
dt=DecisionTreeClassifier(criterion='entropy',splitter='best',max_features=3,max_depth=3)
dt.fit(X_train,Y_train)
Y_pred = dt.predict(X_test)
DT_acc = accuracy_score(Y_test,Y_pred)
```

```
RF=RandomForestClassifier(n_estimators=100,
criterion='entropy',max_depth=3)
RF.fit(X_train,Y_train)
Y_pred = RF.predict(X_test)
RF_acc = accuracy_score(Y_test,Y_pred)
```

```
print("LogisticRegression -> Accuracy -> ",reg_acc)
print("Naïve_Bayes -> Accuracy -> ",bayes_acc)
```

```

print("Decision_Tree -> Accuracy -> ",DT_acc)
print("K-NN -> Accuracy -> ",KNN_acc)
print("SVM -> Accuracy -> ",svc_acc)
print("Random_Forest -> Accuracy -> ",RF_acc)

LogisticRegression -> Accuracy -> 0.631578947368421
Naive_Bayes -> Accuracy -> 0.631578947368421
Decision_Tree -> Accuracy -> 0.9064327485380117
K-NN -> Accuracy -> 0.7660818713450293
SVM -> Accuracy -> 0.631578947368421
Random_Forest -> Accuracy -> 0.935672514619883

RF_roc_auc =metrics.roc_auc_score(Y_test, RF.predict(X_test))


print ("ROC_AUC Score is -> ", RF_roc_auc )

ROC_AUC Score is -> 0.925925925925926

lfpr, ltp, lthresolds = metrics.roc_curve(Y_test,
RF.predict_proba(X_test)[: ,1])
plt.plot(lfpr, ltp, label='AUC (area = %0.2f)' % RF_roc_auc)

plt.plot([0, 1], [0, 1], 'r--')
#plt.xlim([0.0, 1.0])
#plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.show()

```



```
print("F1 Score is : ",f1_score(Y_test, Y_pred))
```

F1 Score is : 0.9105691056910569

```
print(classification_report(Y_test, Y_pred))
```

	precision	recall	f1-score	support
0	0.94	0.96	0.95	108
1	0.93	0.89	0.91	63
accuracy			0.94	171
macro avg	0.94	0.93	0.93	171
weighted avg	0.94	0.94	0.94	171

AS6PR3

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
from sklearn.model_selection import GridSearchCV
```

```
data = pd.read_csv("voice.csv")
data.head()
```

	meanfreq	sd	median	Q25	Q75	IQR
skew \						
0	0.059781	0.064241	0.032027	0.015071	0.090193	0.075122
12.863462						
1	0.066009	0.067310	0.040229	0.019414	0.092666	0.073252
22.423285						
2	0.077316	0.083829	0.036718	0.008701	0.131908	0.123207
30.757155						
3	0.151228	0.072111	0.158011	0.096582	0.207955	0.111374
1.232831						
4	0.135120	0.079146	0.124656	0.078720	0.206045	0.127325
1.101174						

	kurt	sp.ent	sfm	...	centroid	meanfun	minfun
\							
0	274.402906	0.893369	0.491918	...	0.059781	0.084279	0.015702
1	634.613855	0.892193	0.513724	...	0.066009	0.107937	0.015826
2	1024.927705	0.846389	0.478905	...	0.077316	0.098706	0.015656
3	4.177296	0.963322	0.727232	...	0.151228	0.088965	0.017798
4	4.333713	0.971955	0.783568	...	0.135120	0.106398	0.016931

	maxfun	meandom	mindom	maxdom	dfrange	modindx	label
0	0.275862	0.007812	0.007812	0.007812	0.000000	0.000000	male
1	0.250000	0.009014	0.007812	0.054688	0.046875	0.052632	male
2	0.271186	0.007990	0.007812	0.015625	0.007812	0.046512	male
3	0.250000	0.201497	0.007812	0.562500	0.554688	0.247119	male
4	0.266667	0.712812	0.007812	5.484375	5.476562	0.208274	male

[5 rows x 21 columns]

```
data.rename(columns = {'label': 'Gender'}, inplace = True)
data
```

skew \	meanfreq	sd	median	Q25	Q75	IQR
0	0.059781	0.064241	0.032027	0.015071	0.090193	0.075122
12.863462						
1	0.066009	0.067310	0.040229	0.019414	0.092666	0.073252
22.423285						
2	0.077316	0.083829	0.036718	0.008701	0.131908	0.123207
30.757155						
3	0.151228	0.072111	0.158011	0.096582	0.207955	0.111374
1.232831						
4	0.135120	0.079146	0.124656	0.078720	0.206045	0.127325
1.101174						
...
...						
3163	0.131884	0.084734	0.153707	0.049285	0.201144	0.151859
1.762129						
3164	0.116221	0.089221	0.076758	0.042718	0.204911	0.162193
0.693730						
3165	0.142056	0.095798	0.183731	0.033424	0.224360	0.190936
1.876502						
3166	0.143659	0.090628	0.184976	0.043508	0.219943	0.176435
1.591065						
3167	0.165509	0.092884	0.183044	0.070072	0.250827	0.180756
1.705029						

	kurt	sp.ent	sfm	...	centroid	meanfun
minfun \						
0 0.015702	274.402906	0.893369	0.491918	...	0.059781	0.084279
1 0.015826	634.613855	0.892193	0.513724	...	0.066009	0.107937
2 0.015656	1024.927705	0.846389	0.478905	...	0.077316	0.098706
3 0.017798	4.177296	0.963322	0.727232	...	0.151228	0.088965
4 0.016931	4.333713	0.971955	0.783568	...	0.135120	0.106398
...

```

..
3163      6.630383  0.962934  0.763182  ...  0.131884  0.182790
0.083770
3164      2.503954  0.960716  0.709570  ...  0.116221  0.188980
0.034409
3165      6.604509  0.946854  0.654196  ...  0.142056  0.209918
0.039506
3166      5.388298  0.950436  0.675470  ...  0.143659  0.172375
0.034483
3167      5.769115  0.938829  0.601529  ...  0.165509  0.185607
0.062257

```

```

      maxfun  meandom  mindom  maxdom  dfrange  modindx
Gender
0      0.275862  0.007812  0.007812  0.007812  0.000000  0.000000
male
1      0.250000  0.009014  0.007812  0.054688  0.046875  0.052632
male
2      0.271186  0.007990  0.007812  0.015625  0.007812  0.046512
male
3      0.250000  0.201497  0.007812  0.562500  0.554688  0.247119
male
4      0.266667  0.712812  0.007812  5.484375  5.476562  0.208274
male
...      ...      ...      ...      ...      ...      ...
..
3163  0.262295  0.832899  0.007812  4.210938  4.203125  0.161929
female
3164  0.275862  0.909856  0.039062  3.679688  3.640625  0.277897
female
3165  0.275862  0.494271  0.007812  2.937500  2.929688  0.194759
female
3166  0.250000  0.791360  0.007812  3.593750  3.585938  0.311002
female
3167  0.271186  0.227022  0.007812  0.554688  0.546875  0.350000
female

```

[3168 rows x 21 columns]

```

le = LabelEncoder()
data["Gender"] = le.fit_transform(data["Gender"])
data

```

```

      meanfreq      sd      median      Q25      Q75      IQR
skew \
0      0.059781  0.064241  0.032027  0.015071  0.090193  0.075122
12.863462
1      0.066009  0.067310  0.040229  0.019414  0.092666  0.073252
22.423285
2      0.077316  0.083829  0.036718  0.008701  0.131908  0.123207

```



```

30.757155
3      0.151228  0.072111  0.158011  0.096582  0.207955  0.111374
1.232831
4      0.135120  0.079146  0.124656  0.078720  0.206045  0.127325
1.101174
...      ...      ...      ...      ...      ...      ...
...
3163  0.131884  0.084734  0.153707  0.049285  0.201144  0.151859
1.762129
3164  0.116221  0.089221  0.076758  0.042718  0.204911  0.162193
0.693730
3165  0.142056  0.095798  0.183731  0.033424  0.224360  0.190936
1.876502
3166  0.143659  0.090628  0.184976  0.043508  0.219943  0.176435
1.591065
3167  0.165509  0.092884  0.183044  0.070072  0.250827  0.180756
1.705029

```

```

          kurt      sp.ent      sfm      ...      centroid      meanfun
minfun \
0      274.402906  0.893369  0.491918  ...  0.059781  0.084279
0.015702
1      634.613855  0.892193  0.513724  ...  0.066009  0.107937
0.015826
2      1024.927705  0.846389  0.478905  ...  0.077316  0.098706
0.015656
3      4.177296  0.963322  0.727232  ...  0.151228  0.088965
0.017798
4      4.333713  0.971955  0.783568  ...  0.135120  0.106398
0.016931
...      ...      ...      ...      ...      ...      ...
..
3163      6.630383  0.962934  0.763182  ...  0.131884  0.182790
0.083770
3164      2.503954  0.960716  0.709570  ...  0.116221  0.188980
0.034409
3165      6.604509  0.946854  0.654196  ...  0.142056  0.209918
0.039506
3166      5.388298  0.950436  0.675470  ...  0.143659  0.172375
0.034483
3167      5.769115  0.938829  0.601529  ...  0.165509  0.185607
0.062257

```

```

          maxfun      meandom      mindom      maxdom      dfrange      modindx
Gender
0      0.275862  0.007812  0.007812  0.007812  0.000000  0.000000
1
1      0.250000  0.009014  0.007812  0.054688  0.046875  0.052632
1
2      0.271186  0.007990  0.007812  0.015625  0.007812  0.046512

```

```

1
3      0.250000  0.201497  0.007812  0.562500  0.554688  0.247119
1
4      0.266667  0.712812  0.007812  5.484375  5.476562  0.208274
1
...      ...      ...      ...      ...      ...      ...
..
3163  0.262295  0.832899  0.007812  4.210938  4.203125  0.161929
0
3164  0.275862  0.909856  0.039062  3.679688  3.640625  0.277897
0
3165  0.275862  0.494271  0.007812  2.937500  2.929688  0.194759
0
3166  0.250000  0.791360  0.007812  3.593750  3.585938  0.311002
0
3167  0.271186  0.227022  0.007812  0.554688  0.546875  0.350000
0

[3168 rows x 21 columns]

X = data.drop(columns="Gender")
Y= data["Gender"]
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size =
0.3, random_state = 1)

rfc = RandomForestClassifier(n_jobs=-1,max_features=
'sqrt' ,n_estimators=50, oob_score = True)

param_grid = {
    'n_estimators': [50, 300],
    #'max_features': ['auto', 'sqrt', 'log2'],
    #'max_depth' : [4,5,6,7,8],
    'criterion' :['gini', 'entropy']
}

CV_rfc = GridSearchCV(estimator=rfc, param_grid=param_grid, cv= 5)
CV_rfc.fit(X, Y)

GridSearchCV(cv=5,
              estimator=RandomForestClassifier(max_features='sqrt',
                                                n_estimators=50,
n_jobs=-1,
                                                oob_score=True),
              param_grid={'criterion': ['gini', 'entropy'],
                          'n_estimators': [50, 300]})

print("Best Parameter and Estimator are : ")
print (CV_rfc.best_params_)

Best Parameter and Estimator are :
{'criterion': 'gini', 'n_estimators': 50}

```


AS6PR4

```
from sklearn.ensemble import AdaBoostClassifier
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
from sklearn.model_selection import GridSearchCV
```

```
data=pd.read_csv("seeds.csv")
data.head()
```

	Area	Perimeter	Compactness	Kernel.Length	Kernel.Width	\
0	15.26	14.84	0.8710	5.763	3.312	
1	14.88	14.57	0.8811	5.554	3.333	
2	14.29	14.09	0.9050	5.291	3.337	
3	13.84	13.94	0.8955	5.324	3.379	
4	16.14	14.99	0.9034	5.658	3.562	

	Asymmetry.Coeff	Kernel.Groove	Type
0	2.221	5.220	1
1	1.018	4.956	1
2	2.699	4.825	1
3	2.259	4.805	1
4	1.355	5.175	1

```
X = data.drop(columns="Type")
y=data["Type"]
X_train,X_test,y_train,y_test=
train_test_split(X,y,test_size=0.3,random_state=10)
```

```
dt=DecisionTreeClassifier(criterion='entropy',splitter='best',max_dept
h=1)
dt.fit(X_train,y_train)
y_pred = dt.predict(X_test)
DT_acc = accuracy_score(y_test,y_pred)
```

```

gaussian_nb=GaussianNB()
gaussian_nb.fit(X_train,y_train)
y_pred = gaussian_nb.predict(X_test)
bayes_acc = accuracy_score(y_test,y_pred)

```

```

print("Accuracy of Decision Tree Classifier is : ", DT_acc*100)
print("Accuracy of GaussianNB is : ", bayes_acc*100)

```

```

Accuracy of Decision Tree Classifier is : 60.0
Accuracy of GaussianNB is : 91.66666666666666

```

```

ada=AdaBoostClassifier(n_estimators=100,learning_rate=0.6)
ada.fit(X_train,y_train)
ada_pred=ada.predict(X_test)
acc_ada=round(accuracy_score(y_test,ada_pred),3)
print(acc_ada)

```

```

0.85

```

```

ada_gnb=AdaBoostClassifier(base_estimator=GaussianNB(),n_estimators=100,learning_rate=0.5)
ada_gnb.fit(X_train,y_train)
agnb_pred=ada_gnb.predict(X_test)
acc_agnb=round(accuracy_score(y_test,agnb_pred),3)
print(acc_agnb*100)

```

```

91.7

```

```

print ("Unique values are : ", data.Type.unique())

```

```

Unique values are : [1 2 3]

```

```

print("Count of Unique Values are : ", data['Type'].value_counts())

```

```

Count of Unique Values are: 2    68

```

```

1    66

```

```

3    65

```

```

Name: Type, dtype: int64

```

```

print ("Percentage of Unique values are",
data['Type'].value_counts(normalize=True)*100)

```

```

Percentage of Unique values are 2    34.170854

```

```

1    33.165829

```

```

3    32.663317

```

```

Name: Type, dtype: float64

```

```

plt.figure(figsize=(10,2))

```

```

plt.barh(np.arange(4),[DT_acc,bayes_acc,acc_ada,acc_agnb],\

```

```
tick_label=['Decision  
Tree','GaussianNB','AdaBoost','AdaBoost(GaussianNB)']
```

<BarContainer object of 4 artists>

